

# Computer Visions

# Generative Machine Learning Tools

MAS.S68 F'19

Roy Shilkrot, Fluid Interfaces

Class 6: 3D

A dark blue diagonal gradient bar that starts from the bottom left and extends towards the top right, covering the lower half of the slide.

# Today

## Catch up

- HW 3+4 update

## Content

- Representing 3D models
- ML for 3D data
- Generative ML for 3D

## Application

- Sketch to 3D, Photo to 3D
- 3D Latent space learning
- 3D GANs

# Representing 3D Models


Mesh

Graph


Volume

Parametric Design


Generative / Procedural Design




# 3D Model as Mesh




vertices




edges




faces



polygons



surfaces




# 3D Model as Mesh

```
ply
format ascii 1.0
element vertex 867
property float32 x
property float32 y
property float32 z
element face 1704
end_header
0.00472708 0.0012 -0.000833515
0.0048 0.0012 0
0 0 0
0.00451052 0.0012 -0.0016417
0.00415692 0.0012 -0.00240001
0.00367701 0.0012 -0.00308539
0.00308537 0.0012 -0.00367702
0.00239999 0.0012 -0.00415693
0.00164168 0.0012 -0.00451053
3 0 1 2
3 3 0 2
3 4 3 2
3 5 4 2
3 6 5 2
3 7 6 2
3 8 7 2
3 9 8 2
3 10 9 2
```




## Face-Vertex Meshes

	Face List
f0	v0 v4 v5
f1	v0 v5 v1
f2	v1 v5 v6
f3	v1 v6 v2
f4	v2 v6 v7
f5	v2 v7 v3
f6	v3 v7 v4
f7	v3 v4 v0
f8	v8 v5 v4
f9	v8 v6 v5
f10	v8 v7 v6
f11	v8 v4 v7
f12	v9 v5 v4
f13	v9 v6 v5
f14	v9 v7 v6
f15	v9 v4 v7



	Vertex List
v0	0,0,0 f0 f1 f12 f15 f7
v1	1,0,0 f2 f3 f13 f12 f1
v2	1,1,0 f4 f5 f14 f13 f3
v3	0,1,0 f6 f7 f15 f14 f5
v4	0,0,1 f6 f7 f0 f8 f11
v5	1,0,1 f0 f1 f2 f9 f8
v6	1,1,1 f2 f3 f4 f10 f9
v7	0,1,1 f4 f5 f6 f11 f10
v8	.5,.5,0 f8 f9 f10 f11
v9	.5,.5,1 f12 f13 f14 f15

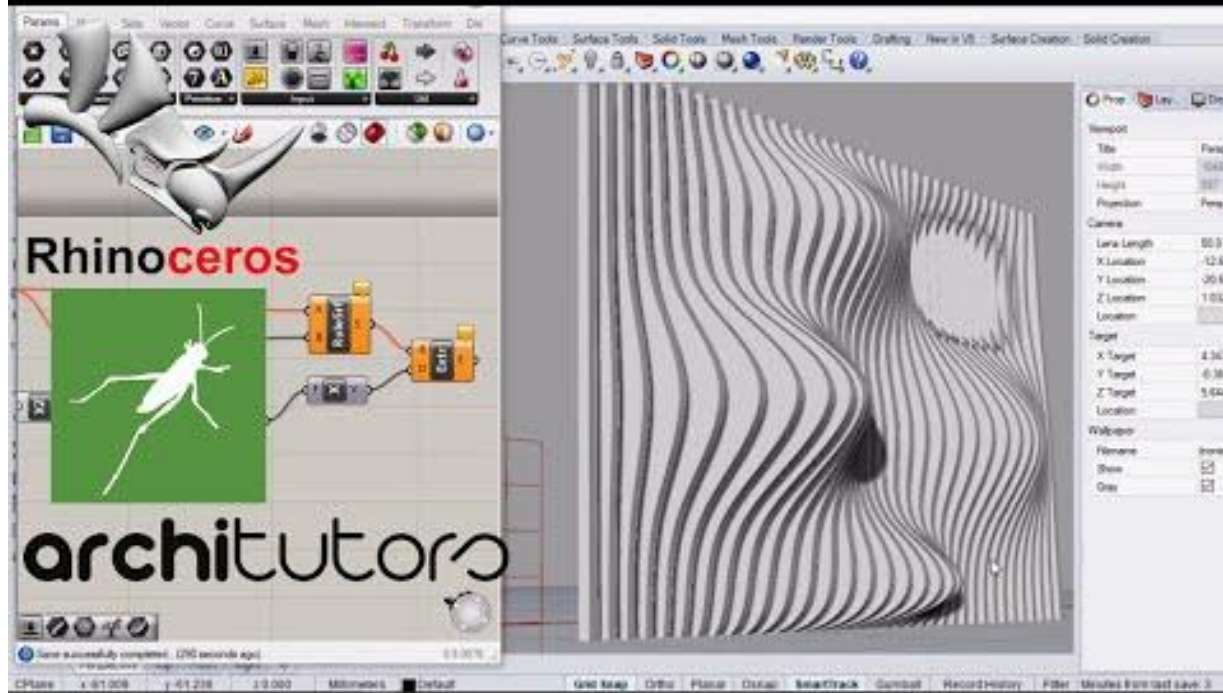


# 3D Model as Solid



# Parametric 3D Modeling






Parametric Design







Parametric Design (Architecture)



Generative Design (Architecture)





Generative Design



Procedural Design


# 3D Model as Volume

Voxel



Mesh to Volume (Voxelization)










Volumetric Modeling

# 3D Volume: Model as Distance Function

www.scratchapixel.com



## DeepSDF: Learning Continuous Signed Distance Functions for Shape Representation



Learned Chair Shape Space




Learned Car Shape Space

Example Linear Interpolation through DeepSDF latent shape space trained on 4116 chairs and 662 car models




# Deep Learning for 3D Models

3D Deep Learning Tasks  
3D Mesh as Graph  
Graphs Learning




(1)



(1')

## OctNet

· Pooling





Input



Fan et al. (2017)



Predicting 3D Volume from 2D Images

# Edge Centric 3D Model

Face List


f0	4 8 9
f1	0 10 9
f2	5 10 11
f3	1 12 11
f4	6 12 13
f5	2 14 13
f6	7 14 15
f7	3 8 15
f8	4 16 19
f9	5 17 16
f10	6 18 17
f11	7 19 18
f12	0 23 20
f13	1 20 21
f14	2 21 22
f15	3 22 23

Edge List



e0	v0 v1	f1 f12	9 23 10 20
e1	v1 v2	f3 f13	11 20 12 21
e2	v2 v3	f5 f14	13 21 14 22
e3	v3 v0	f7 f15	15 22 8 23
e4	v4 v5	f0 f8	19 8 16 9
e5	v5 v6	f2 f9	16 10 17 11
e6	v6 v7	f4 f10	17 12 18 13
e7	v7 v4	f6 f11	18 14 19 15
e8	v0 v4	f7 f0	3 9 7 4
e9	v0 v5	f0 f1	8 0 4 10
e10	v1 v5	f1 f2	0 11 9 5
e11	v1 v6	f2 f3	10 1 5 12
e12	v2 v6	f3 f4	1 13 11 6
e13	v2 v7	f4 f5	12 2 6 14
e14	v3 v7	f5 f6	2 15 13 7
e15	v3 v4	f6 f7	14 3 7 15
e16	v5 v8	f8 f9	4 5 19 17
e17	v6 v8	f9 f10	5 6 16 18
e18	v7 v8	f10 f11	6 7 17 19
e19	v4 v8	f11 f8	7 4 18 16
e20	v1 v9	f12 f13	0 1 23 21
e21	v2 v9	f13 f14	1 2 20 22
e22	v3 v9	f14 f15	2 3 21 23
e23	v0 v9	f15 f12	3 0 22 20

Vertex List

v0	0,0,0	8 9 0 23 3
v1	1,0,0	10 11 1 20 0
v2	1,1,0	12 13 2 21 1
v3	0,1,0	14 15 3 22 2
v4	0,0,1	8 15 7 19 4
v5	1,0,1	10 9 4 16 5
v6	1,1,1	12 11 5 17 6
v7	0,1,1	14 13 6 18 7
v8	5,5,0	16 17 18 19
v9	5,5,1	20 21 22 23



Winged-Edge Meshes



# Graphs for the Deep Learning Era

**Graph Neural Network**  $\mathbf{h}_v = f(\mathbf{x}_v, \mathbf{x}_{co[v]}, \mathbf{h}_{ne[v]}, \mathbf{x}_{ne[v]})$

Hidden state:  $\mathbf{H}^{t+1} = F(\mathbf{H}^t, \mathbf{X})$

$$\mathbf{o}_v = g(\mathbf{h}_v, \mathbf{x}_v)$$

Output:

$$loss = \sum_{i=1}^p (\mathbf{t}_i - \mathbf{o}_i)$$

Node  
labeling

Loss:


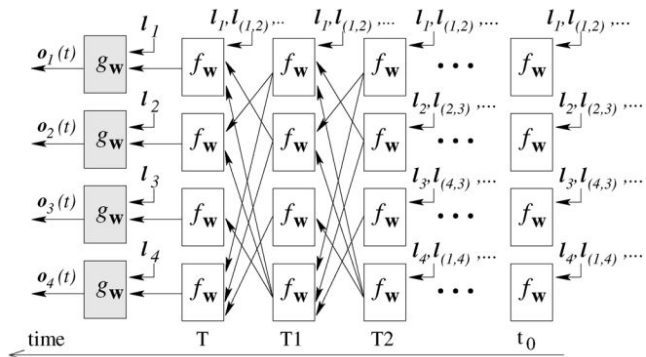
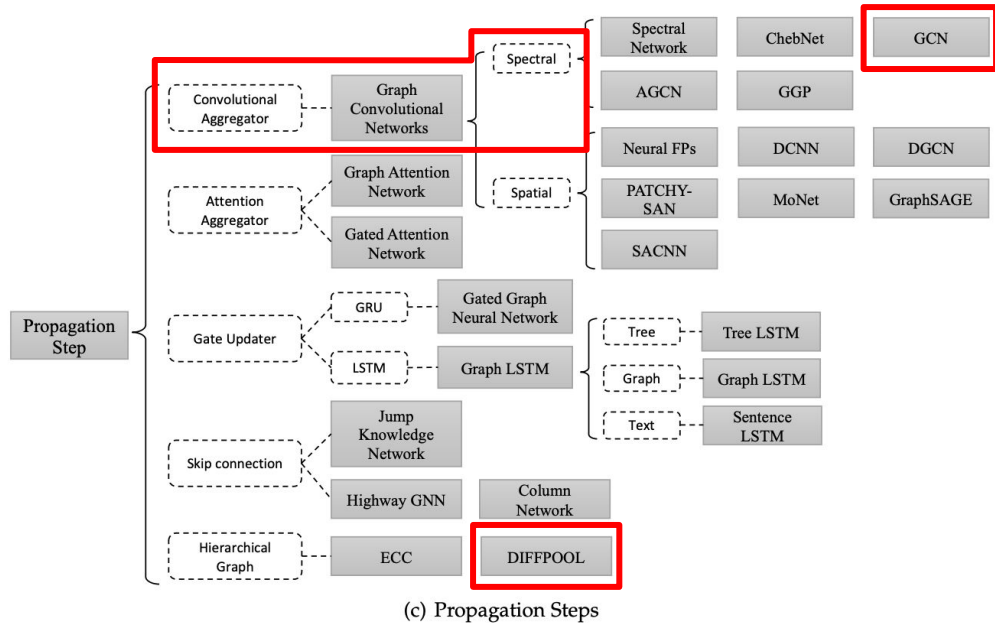
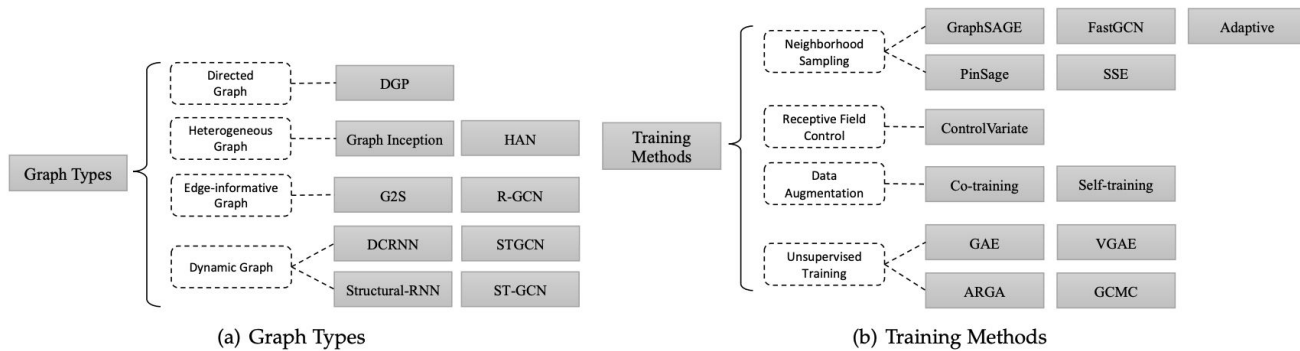


Fig. 2. Graph and the neighborhood of a node. The state  $x_1$  of the node 1 depends on the information contained in its neighborhood.

[Scarselli '09]



Us, today.

Fig. 2. An overview of variants of graph neural networks.

[Zhou '18]

# Graph Convolutional Networks

Popularized by: [Kipf and Welling '16](#), ICLR '17

But it existed long before...


Graph Convolution Layer:

$$H^{(l+1)} = \sigma \left( \overset{\text{hidden}}{\tilde{D}}^{-\frac{1}{2}} \overset{\text{adjacency}}{\tilde{A}} \tilde{D}^{-\frac{1}{2}} \overset{\text{weights}}{H^{(l)}} W^{(l)} \right)$$

$$\tilde{A} = A + I_N \quad \tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$$

Learn linear transformation of node features (embeddings), add non-linear activation.

Why “**Convolution**”?



(a) Graph Convolutional Network



# Graph Convolutions

Symmetric normalized graph **Laplacian**:

$$\mathbf{L} = \mathbf{I}_n - \mathbf{D}^{-\frac{1}{2}} \mathbf{A} \mathbf{D}^{-\frac{1}{2}}$$

A (adjacency) 1s or 0s, its diagonal is all 0s  
D degree matrix.

*What is it useful for??*

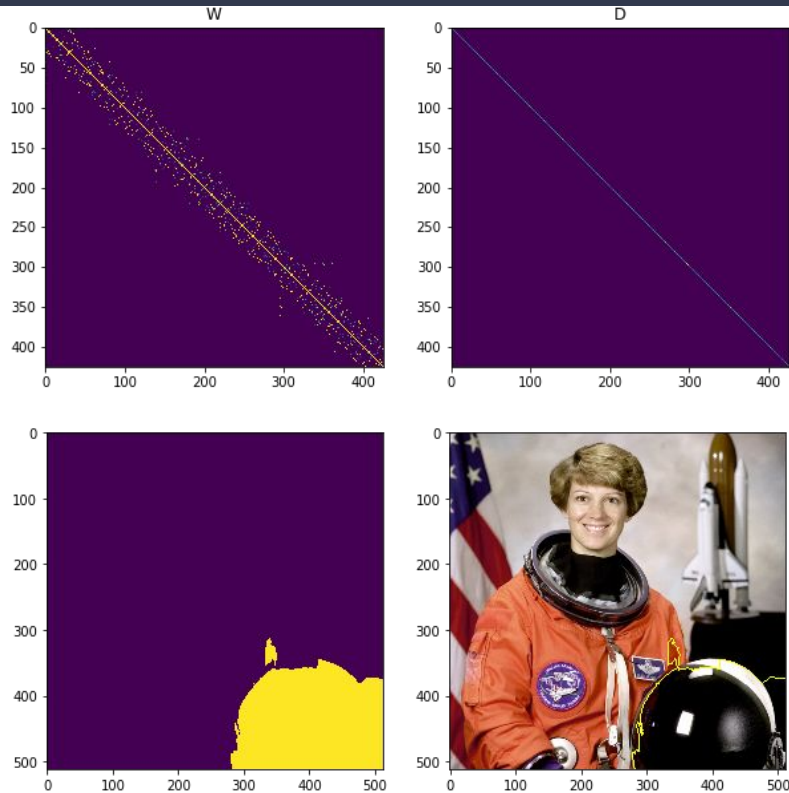
Well, it's actually useful for Segmentation!

[Shi, Jianbo, Malik '00] **Normalized Cuts**

Taking the second smallest eigenvector from:

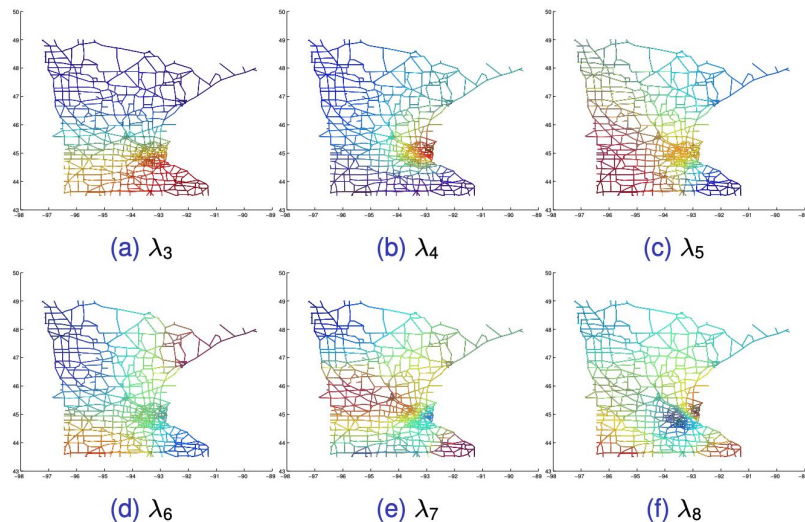
$$\mathbf{L} = \mathbf{U} \mathbf{\Lambda} \mathbf{U}^T \quad \mathbf{L}_{ii} = \lambda_i$$

(In norm-cuts the non-sym Lap. is  $\mathbf{L} = \mathbf{D} - \mathbf{A}$ )

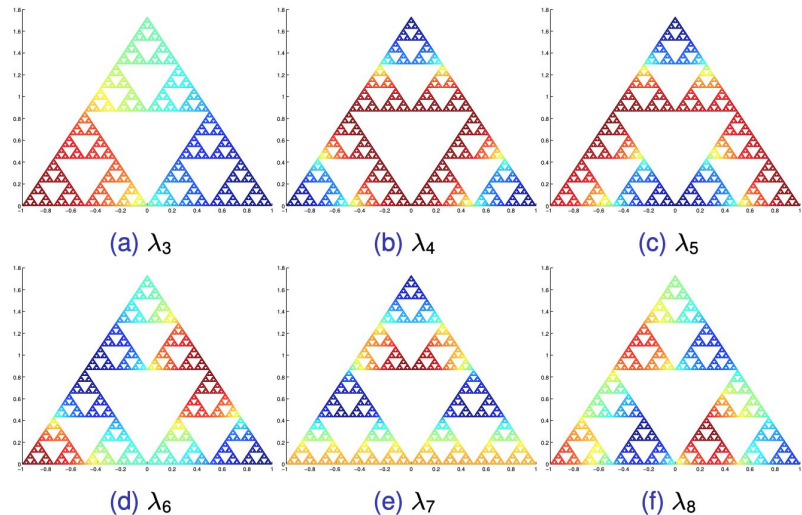


# Graph Convolutions

Graph Laplacian spectral analysis (e.g. eigen decomposition) exposes structure.



**Figure :** Eigenfunctions corresponding to the first six nonzero eigenvalues. Minnesota road graph (2642 vertices)



**Figure :** Eigenfunctions corresponding to the first six nonzero eigenvalues. Level-8 graph approximation to Sierpinski gasket (9843 vertices)

# Graph Convolutions

Graph Fourier Transform: (spectral / eigen linear basis transform)

[Wu '19]

$$\begin{aligned}\mathbf{L} &= \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T \\ \mathcal{F}(\mathbf{x}) &= \mathbf{U}^T \mathbf{x} \quad \mathcal{F}^{-1}(\hat{\mathbf{x}}) = \mathbf{U}\hat{\mathbf{x}} \\ \mathbf{x} &= \sum_i \hat{\mathbf{x}}_i \mathbf{u}_i\end{aligned}$$

Graph filter, via convolution theorem:

$$\begin{aligned}\mathbf{g} &\in \mathbf{R}^N \\ \mathbf{x} *_G \mathbf{g} &= \mathcal{F}^{-1}(\mathcal{F}(\mathbf{x}) \odot \mathcal{F}(\mathbf{g})) \\ &= \mathbf{U}(\mathbf{U}^T \mathbf{x} \odot \mathbf{U}^T \mathbf{g})\end{aligned}$$

Linear simplification:

$$\mathbf{g}_\theta = \text{diag}(\mathbf{U}^T \mathbf{g}) \longrightarrow \mathbf{x} *_G \mathbf{g}_\theta = \mathbf{U} \mathbf{g}_\theta \mathbf{U}^T \mathbf{x}$$

# Graph Convolutions

Generalized convolution:

$$\mathbf{x} *_G \mathbf{g}_\theta = \mathbf{U} \mathbf{g}_\theta \mathbf{U}^T \mathbf{x}$$

[[Kipf '16](#)]

Solving for  $\mathbf{g}_\theta$  requires K-th order polynomials in the Laplacian, and depends on nodes that are K steps away from the central node (K-th-order neighborhood).

[Kipf & Welling '16] suggest taking  $K = 1$ , and end up with the simpler linear transformation:

$$\mathbf{g}_\theta \star x \approx \theta \left( I_N + D^{-\frac{1}{2}} A D^{-\frac{1}{2}} \right) x$$

$$\mathbf{Z} = \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} \mathbf{X} \Theta \quad \mathbf{X} \in \mathbb{R}^{N \times C}$$


It is **differentiable** in  $\Theta$ , and **fast to compute** both FF and Backprop.

And the beauty is that it follows CNN paradigm: **Stack layers!**

# Receptive Field Propagation

By stacking GC layers we increase the receptive field of higher-level neurons.

However there's one **big** problem ---



# Receptive Field Propagation


By stacking GC layers we increase the receptive field of higher-level neurons.


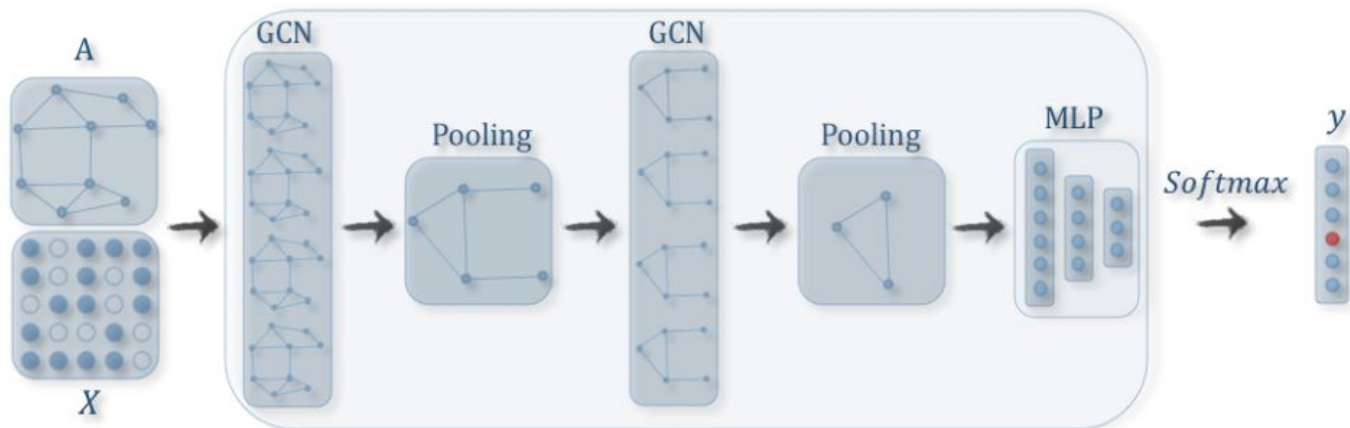
However there's one **big** problem ---

Graph *layout* stays the same layer-to-layer.  
We have the *same amount of features* to learn in each layer === it doesn't scale.

It's essentially a MLP.

(Quick to overfit, bloated, data hungry models)



**Node  
classification****Graph  
classification**



# DiffPool and conv-pool




Figure 1: High-level illustration of our proposed method DIFFPOOL. At each hierarchical layer, we run a GNN model to obtain embeddings of nodes. We then use these learned embeddings to cluster nodes together and run another GNN layer on this coarsened graph. This whole process is repeated for  $L$  layers and we use the final output representation to classify the graph.

[Ying '18]



# Learning How to “Pool” a Graph

[Ying '18]

Select the nodes (via indicator matrix) that will go to the next cluster.

Learn this as a *parametric assignment*.

E.g. the assignment is dependant on the node features + neighbor features.

$$A^{(i+1)} = S^{(i)\top} A^{(i)} S^{(i)}$$

$$S^{(i)} = \text{softmax} \left( \text{GNN}_{\text{pool}} \left( A^{(i)}, F^{(i)} \right) \right)$$

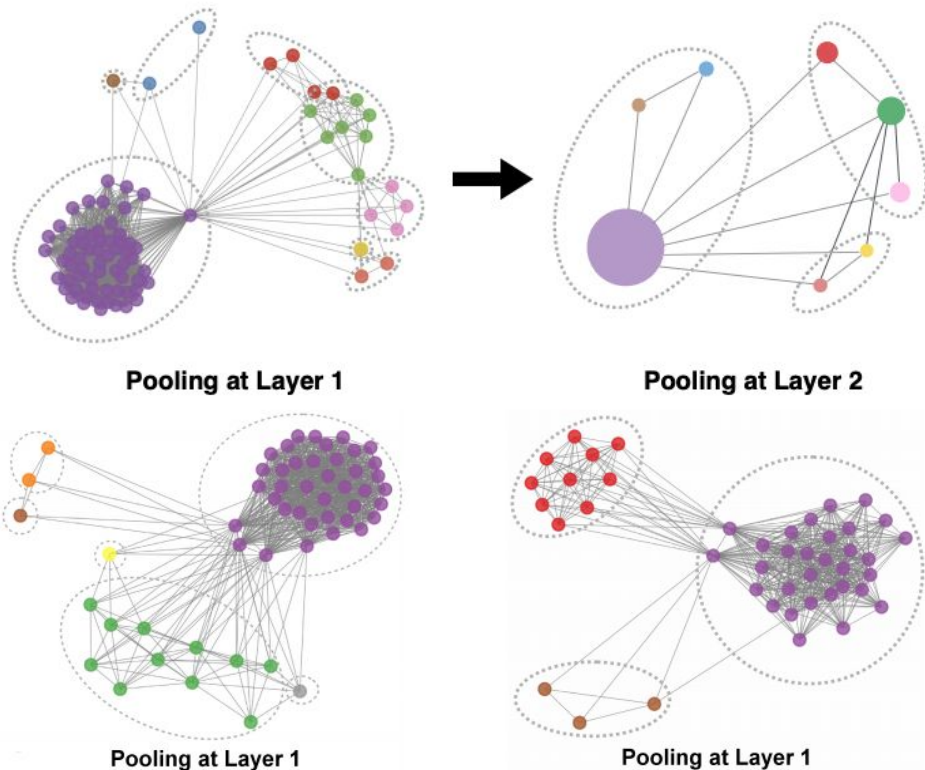
[Kipf et al]


Encourage neighbors to *stick together*:

$$L_{\text{LP}} = \|A^{(l)}, S^{(l)} S^{(l)\top}\|_F$$

Enforce *singular assignment* via entropy:

$$L_{\text{E}} = \frac{1}{n} \sum_{i=1}^n H(S_i)$$





Graph Convolutions for 3D Mesh Segmentation

# Generative 3D Models

3DGAN

Caricature to 3D

Image to 3D

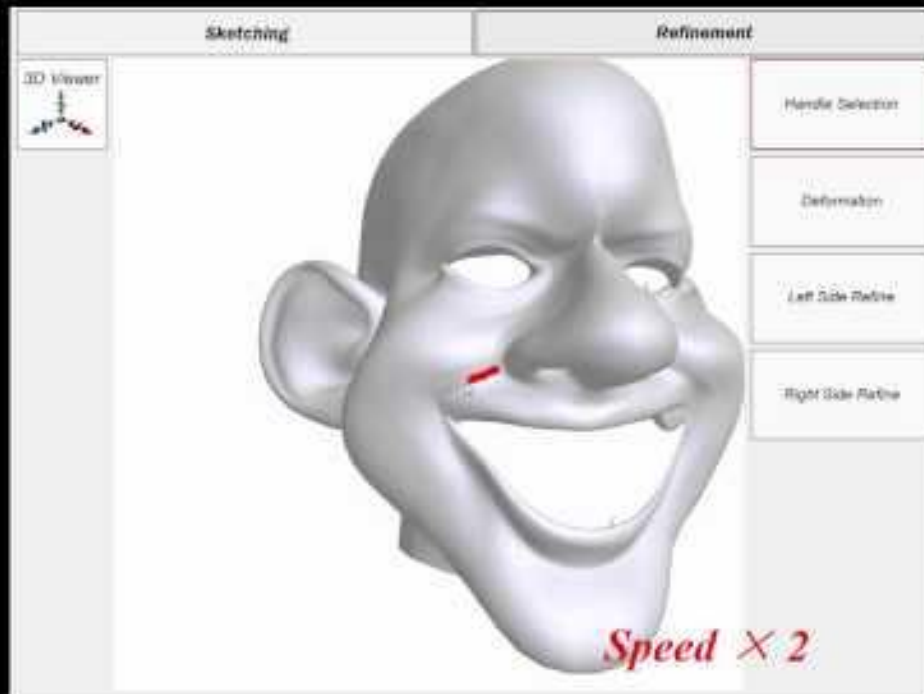
HoloGAN



## Our Synthesized 3D Shapes



In this work, we build a model to generate 3D shapes from latent vectors.




Sketch-to-3D



3D from Image

Given a mesh of a car, we first compute depth maps from different viewpoints.



# HoloGAN

Unsupervised learning of 3D representations from natural images

Thu Nguyen-Phuoc · Chuan Li · Lucas Theis · Christian Richardt · Yong-Liang Yang

ICCV 2019





